

Generative AI & Prompt Engineering — Technical Deliverable

Backend Node.js Track | Information Technology Institute (ITI)

1. Framework Comparison & Evaluation (LangChain.js vs. LlamaIndex.TS)

When building Retrieval-Augmented Generation (RAG) pipelines in TypeScript/Node.js environments, both **LangChain.js** and **LlamaIndex.TS** offer powerful abstractions. Below is an in-depth evaluation based on core architectural capabilities.

A. Ease of Integration

- **LangChain.js:** Follows a highly modular, chain-based approach using the *LangChain Expression Language (LCEL)*. While extremely flexible, it has a steeper learning curve because developers must explicitly chain models, prompts, output parsers, and retrievers together.
- **LlamaIndex.TS:** Adopts a data-first approach with high-level abstractions designed specifically for ingestion, indexing, and RAG. It allows developers to build a working, production-grade RAG pipeline with significantly less boilerplate code.

B. Flexibility with Vector Databases

Both frameworks seamlessly integrate with industry-standard vector databases (**Pinecone**, **ChromaDB**, **Qdrant**):

- **LangChain.js:** Excellent for complex workflows requiring hybrid search, granular metadata filtering across various stores, and custom retriever implementations via [@langchain/community](#).
- **LlamaIndex.TS:** Specialized in optimized text chunking, node parsing, and hierarchical indexing structures (e.g., **VectorStoreIndex**, **SummaryIndex**), providing superior out-of-the-box data retrieval accuracy.

C. Chat History & Memory Management

- **LangChain.js:** Provides comprehensive built-in memory utilities (**BufferMemory**, **ConversationSummaryWindowMemory**) with direct persistence options for Redis, MongoDB, and PostgreSQL.
- **LlamaIndex.TS:** Offers robust state handling via **ContextChatEngine**, automatically combining retrieved documents with active conversation context.

Evaluation Matrix

Criterion	LangChain.js	LlamaIndex.TS
Primary Focus		

Criterion	LangChain.js	LlamaIndex.TS
	General AI Applications, Multi-step Chains & Autonomous Agents	Data Indexing, Structuring & Advanced RAG Pipelines
Learning Curve	Moderate to High	Low to Moderate
RAG Abstractions	Modular & Low-level (LCEL)	High-level & Specialized
Memory Management	Extensive (Buffer, Summary, Hybrid Stores)	Engine-level Contextual Memory

2. Concrete Code Implementations (TypeScript)

A. RAG Pipeline Implementation using LangChain.js & Vector DB

The following example demonstrates how to set up an end-to-end RAG pipeline using LangChain.js with Pinecone and OpenAI:

```
import { OpenAIEmbeddings, ChatOpenAI } from "@langchain/openai";
import { PineconeStore } from "@langchain/pinecone";
import { Pinecone } from "@pinecone-database/pinecone";
import { StringOutputParser } from "@langchain/core/output_parsers";
import { RunnableSequence, RunnablePassthrough } from "@langchain/core/runnables";
import { PromptTemplate } from "@langchain/core/prompts";

export async function runLangChainRAG(userQuery: string): Promise<string> {
  // 1. Initialize Pinecone Client
  const pinecone = new Pinecone({ apiKey: process.env.PINECONE_API_KEY! });
  const pineconeIndex = pinecone.Index(process.env.PINECONE_INDEX_NAME!);

  // 2. Initialize Embeddings and Vector Store Retriever
  const embeddings = new OpenAIEmbeddings({ modelName: "text-embedding-3-small" });
  const vectorStore = await PineconeStore.fromExistingIndex(embeddings, {
    pineconeIndex,
  });
  const retriever = vectorStore.asRetriever({ k: 3 });

  // 3. Define Prompt Template
  const prompt = PromptTemplate.fromTemplate(`
    Answer the question based ONLY on the following context:
    {context}

    Question: {question}
    Answer:
  `);
}
```

```

// 4. Initialize LLM Model
const model = new ChatOpenAI({ modelName: "gpt-4o", temperature: 0 });

// 5. Construct LCEL Chain
const chain = RunnableSequence.from([
  {
    context: retriever.pipe((docs) => docs.map((d) => d.pageContent).join("
")),
    question: new RunnablePassthrough(),
  },
  prompt,
  model,
  new StringOutputParser(),
]);

// 6. Execute Chain
const response = await chain.invoke(userQuery);
return response;
}

```

B. Practical Implementation of Semantic Caching using Redis

Semantic caching stores vector embeddings of queries in Redis to match semantically identical questions, drastically reducing LLM token costs and latency.

```

import { createClient } from "redis";
import { OpenAIEmbeddings } from "@langchain/openai";

export class SemanticCache {
  private redisClient;
  private embeddings: OpenAIEmbeddings;
  private SIMILARITY_THRESHOLD = 0.92; // Cosine similarity threshold (92%)

  constructor() {
    this.redisClient = createClient({ url: process.env.REDIS_URL });
    this.redisClient.connect();
    this.embeddings = new OpenAIEmbeddings({ modelName: "text-embedding-3-small" });
  }

  // Calculate Cosine Similarity between two vectors
  private cosineSimilarity(vecA: number[], vecB: number[]): number {
    const dotProduct = vecA.reduce((sum, a, idx) => sum + a * vecB[idx], 0);
    const normA = Math.sqrt(vecA.reduce((sum, a) => sum + a * a, 0));
    const normB = Math.sqrt(vecB.reduce((sum, b) => sum + b * b, 0));
    return dotProduct / (normA * normB);
  }

  async getCachedResponse(prompt: string): Promise<string | null> {
    const queryVector = await this.embeddings.embedQuery(prompt);
    const keys = await this.redisClient.keys("cache:*");
  }
}

```

```

    for (const key of keys) {
      const cachedData = await this.redisClient.hGetAll(key);
      if (cachedData.vector) {
        const storedVector = JSON.parse(cachedData.vector);
        const similarity = this.cosineSimilarity(queryVector, storedVector);

        if (similarity >= this.SIMILARITY_THRESHOLD) {
          console.log(`[Semantic Cache HIT] Similarity: ${similarity * 100}.toFixed(2)}%`);
          return cachedData.response;
        }
      }
    }
    console.log("[Semantic Cache MISS]");
    return null;
  }

  async setCachedResponse(prompt: string, response: string): Promise<void> {
    const queryVector = await this.embeddings.embedQuery(prompt);
    const cacheKey = `cache:${Date.now()}`;

    await this.redisClient.hSet(cacheKey, {
      prompt,
      response,
      vector: JSON.stringify(queryVector),
    });
  }
}

```

3. Cost Optimization via Semantic Caching

A. Architectural Concept

Traditional HTTP caching relies on exact string matches. If User A asks *"What is the refund policy?"* and User B asks *"How can I get my money back?"*, traditional caching fails.

Semantic Caching converts incoming queries into vector embeddings and performs a nearest-neighbor vector search in Redis. If a previous query shares high semantic similarity ($\geq 92\%$), the system returns the cached answer instantly, bypassing the LLM call entirely.

B. Key Business & Performance Benefits

1. **Significant API Cost Reduction:** Eliminates duplicate calls to expensive APIs (e.g., GPT-4o), saving up to 40-60% on token consumption in high-volume production environments.
2. **Sub-Millisecond Query Latency:** LLM response generation typically takes 1,000ms – 3,000ms. Returning a cached result from Redis takes **< 10ms**, dramatically improving user experience.

3. **Rate-Limit Mitigation:** Protects your backend infrastructure from hitting provider rate limits during peak usage spikes.